

Spectral and neural gas clustering techniques for software defect prediction: performance evaluation with different feature selection methods

Gianluca Bertaccini

Master's degree in Statistical Sciences (Data Science)
University of Bologna

gianluca.bertaccini@studio.unibo.it

March 29, 2022

Software defect prediction

- Predicting software defects is essential for software developers: if they are able to **identify in advance defective instances of code**, they can spend their time and resources fixing them rather than the clean ones.
- Supervised and unsupervised models can correctly identify software defects by analyzing datasets of **software metrics**, directly computed from the source code.
- In general, software defect prediction makes the process of software development more efficient.

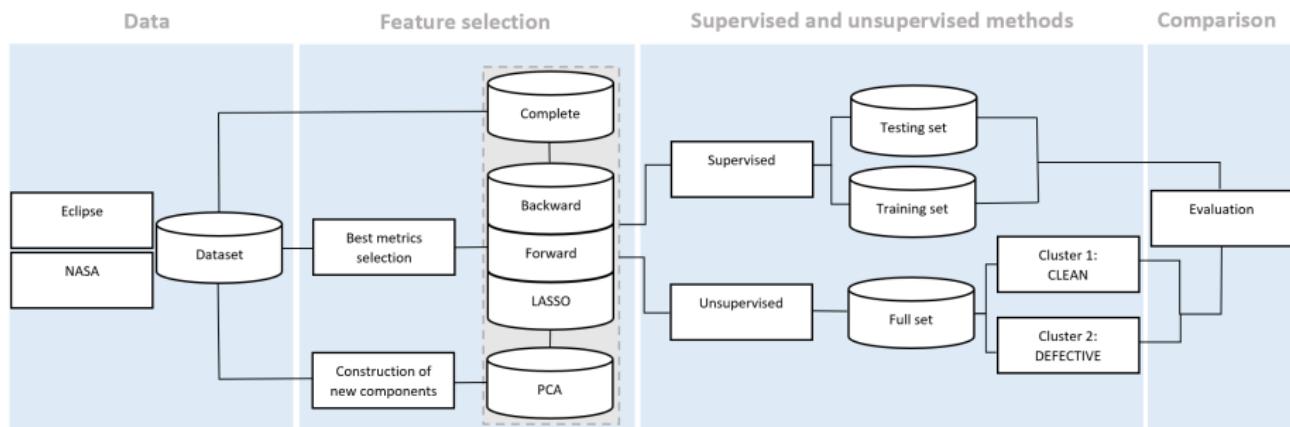
Research questions

These are the research questions of software defect prediction that are considered in this project:

- **RQ1.** Do current unsupervised methods perform better or worse than the supervised ones?
- **RQ2.** Which feature selection method is the best for supervised methods?
- **RQ3.** Which feature selection method is the best for unsupervised methods?

Methodology

Methodology followed to answer the research questions:



Data

To each project corresponds a dataset. There are in total:

- 5 Eclipse datasets
- 12 NASA datasets

Eclipse projects	Description
Eclipse JDT Core	IDE for software developers
Equinox Framework	for development of modular software programs
Apache Lucene	indexing and search features
Mylyn	task management tools
Eclipse PDE UI	plug-in development environment to develop Eclipse features
NASA projects	Description
CM1	software for a spacecraft instrument
JM1	simulate ground predictions
KC1, KC3	storage management software
MC1, MC2, MW2	no information available
PC1, PC2, PC3, PC4, PC5	software satellite

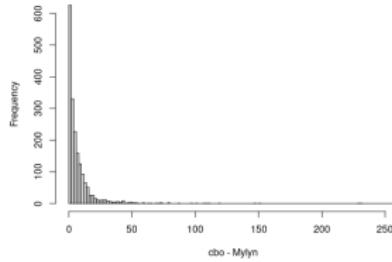
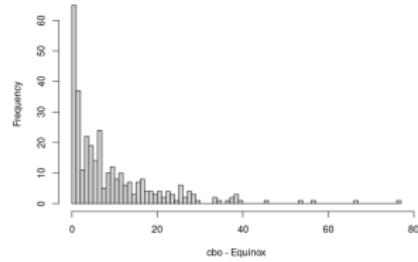
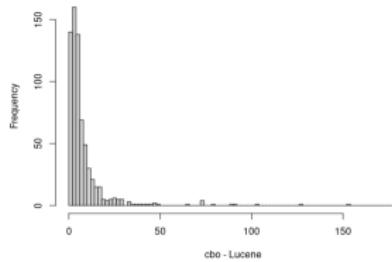
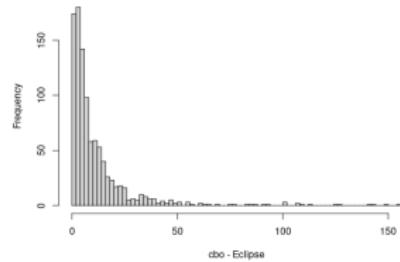
Data: Eclipse

Eclipse datasets contain 18 numeric variables:

- **6 CK** software metrics: they describe the main characteristics of a class (or different instance) of code.
 - **cbo**: coupling between objects
 - **dit**: depth of inheritance tree
 - **lcom**: lack of cohesion methods
 - **noc**: number of children
 - **rfc**: response for a class
 - **wmc**: weighted method count
- **11 object-oriented** software metrics. They describe further characteristics of classes of code, such as number of lines, number of methods or number of attributes.
- **1 binary response** variable, labeling the instance as defective or not (1 or 0).

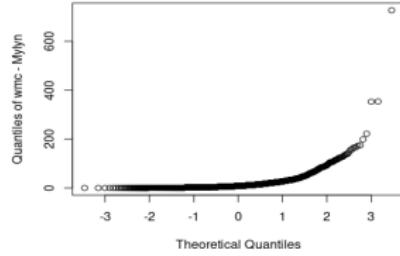
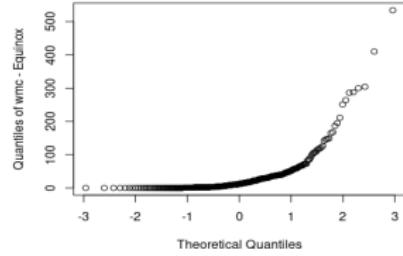
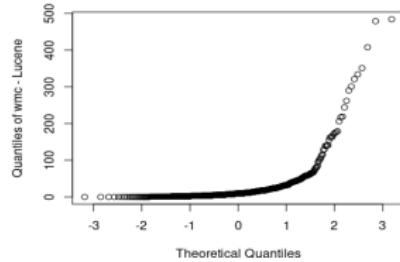
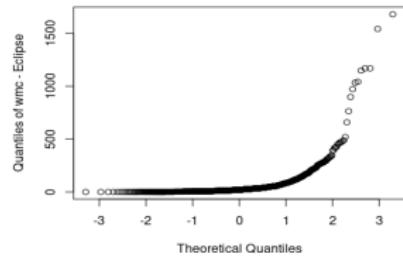
Data: Eclipse, histograms for cbo variable

The data is not normally distributed. For example, the histograms of cbo variable are highly skewed.



Data: Eclipse, quantile plots for wmc variable

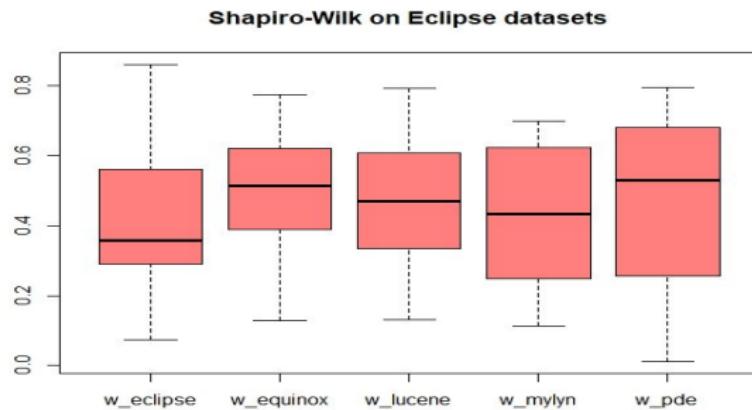
The data is not normally distributed. For example, the quantile plots of the wmc variable show a strong deviation from normality.



Data: Eclipse, Shapiro-Wilk tests

The average Shapiro-Wilk test for every dataset is in the interval between 0.4 and 0.6, which means the variables are not normally distributed. For example, these are the variables lcom and nopa:

Variables	W_Eclipse	W_Equinox	W_Lucene	W_Mylyn	W_PDE
lcom	0.07	0.37	0.13	0.21	0.37
nopa	0.14	0.13	0.47	0.11	0.01



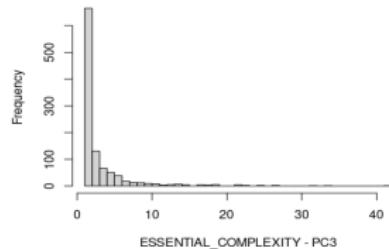
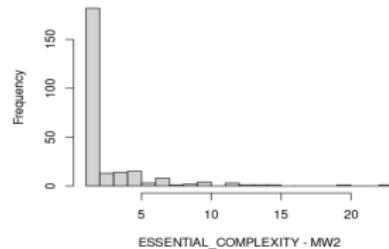
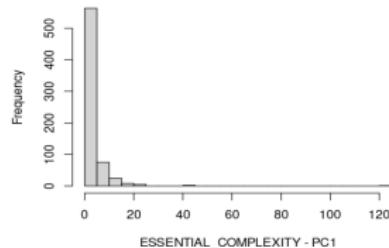
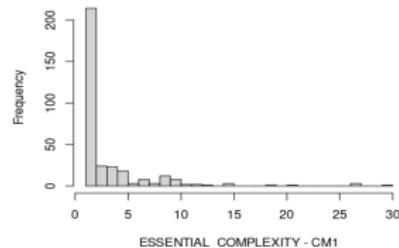
Data: NASA

NASA datasets contain different number of variables, from 22 to 40 variables in total (depending on the dataset). These variables are all numeric and belong to different categories:

- **Lines of code metrics:** counting the number of lines of code, comments, percentages.
- **McCabe's complexity metrics:** metrics computed starting from the McCabe's Cyclomatic Complexity (number of conditional branches in the control-flow graph of an instance of code), described as $M = E - N + 2P$ where E is the number of edges, N is the number of nodes and P is the number of connected components.
- **Halstead complexity metrics:** different complexity metrics that are computed starting from the number of operands and operators.
- **Other complexity metrics:** complexity metrics that are not based directly on McCabe's or Halstead's complexity.
- **Density metrics:** complexity metrics with respect to the number of total paths in an instance of code.
- **1 binary response variable**, labeling the instance as defective or not.

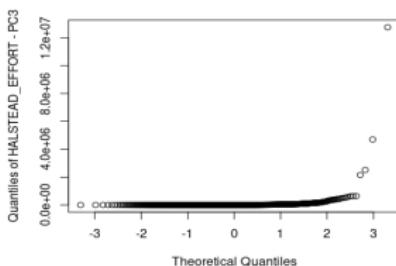
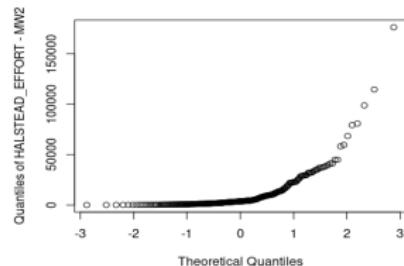
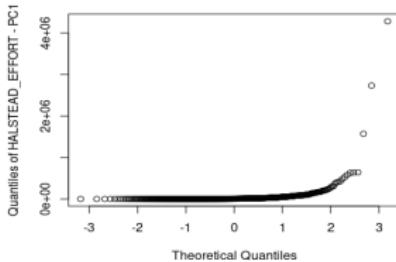
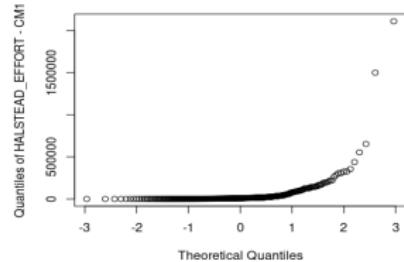
Data: NASA, histograms for ESSENTIAL_COMPLEXITY

The data is not normally distributed. For example, the histograms of ESSENTIAL_COMPLEXITY variable are highly skewed.



Data: NASA, quantile plots for HALSTEAD_EFFORT

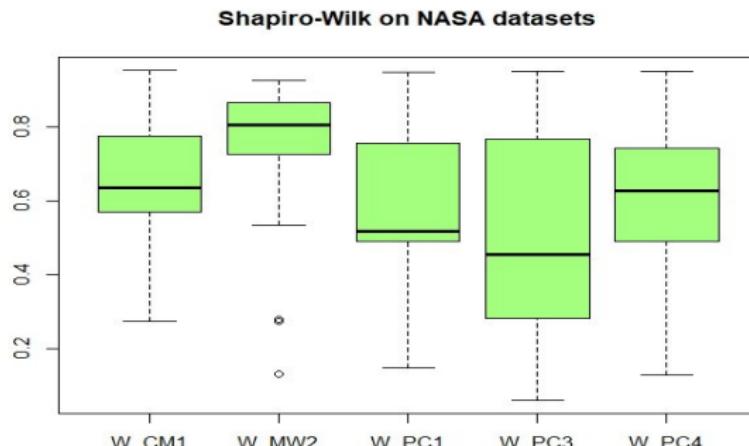
The data is not normally distributed. For example, the quantile plots of the HALSTEAD_EFFORT variable show a strong deviation from normality.



Data: NASA, Shapiro-Wilk tests

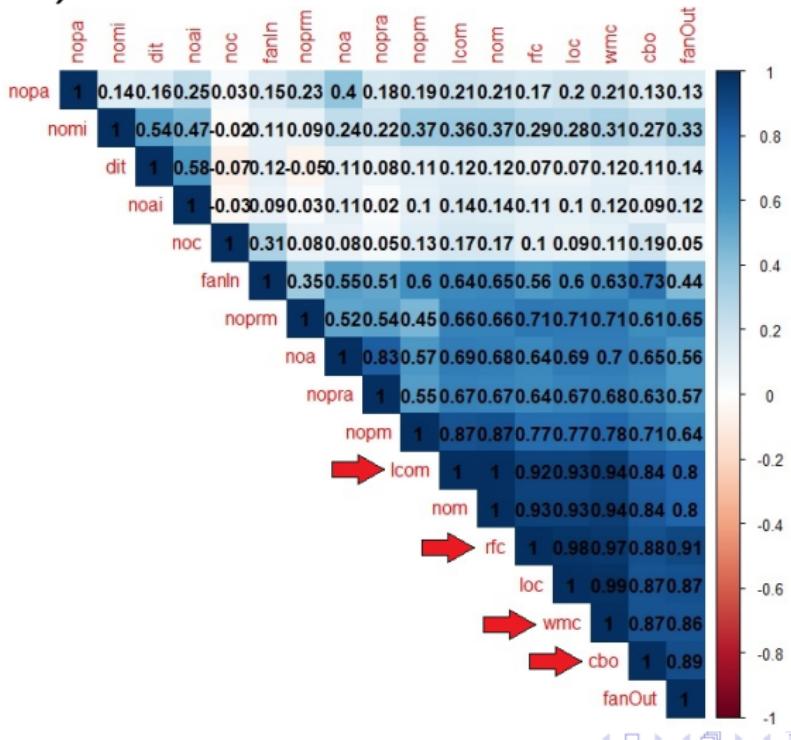
The average Shapiro-Wilk test for NASA datasets is a bit higher than for the Eclipse ones. It reaches up to 0.8. Most variables are still not normally distributed, except for:

Variables	W_Eclipse	W_Equinox	W_Lucene	W_Mylyn	W_PDE
CYCLOMATIC_DENSITY	0.91	0.93	0.95	0.95	0.88
NORMALIZED_CYCLOMATIC_COMPLEXITY	0.86	0.92	0.94	0.93	0.85
NUM_UNIQUE_OPERATORS	0.86	0.90	0.82	0.90	0.95
PERCENT_COMMENTS	0.95	0.91	0.82	0.83	0.85



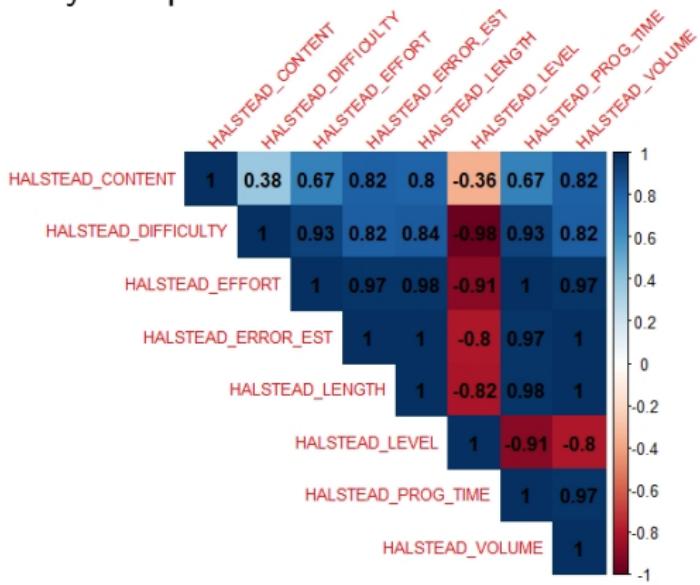
Correlations: Eclipse

Eclipse datasets have high correlations for the **CK** metrics (dit, noc, rfc, wmc, cbo, lcom).



Correlations: NASA

NASA datasets have high correlations for the **Halstead** metrics. This can influence negatively the performance of machine learning methods.

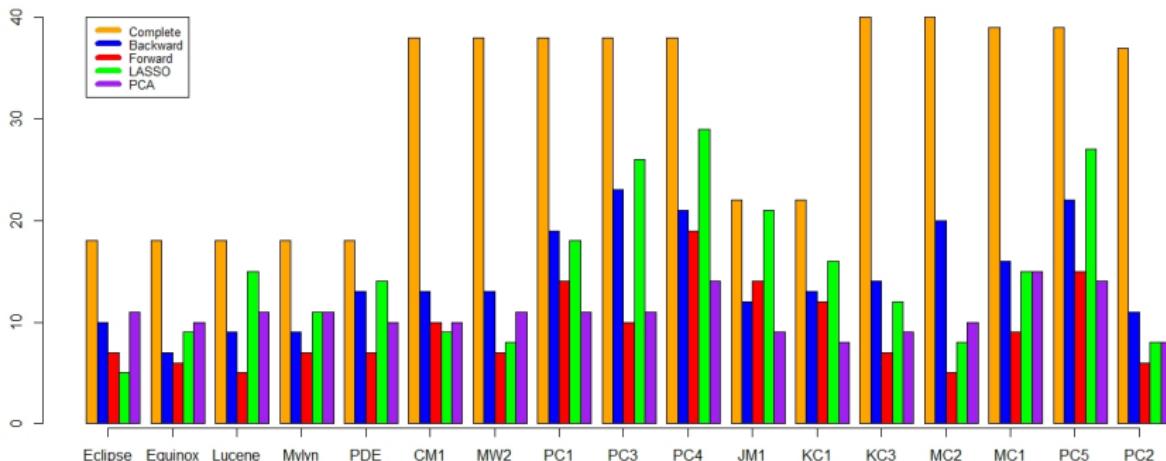


Feature selection

- To reduce the correlations between software metrics, different approaches of feature selection (or feature construction) are tested with each machine learning method.
- Choosing a smaller subset of features can reduce correlations and improve the algorithms performance.
- The approaches are:
 - **Backward** Elimination
 - **Forward** Selection
 - Least Absolute Shrinkage and Selection Operator (**LASSO**)
 - Principal Components Analysis (**PCA**)

Feature selection

Number of features selected by different feature selection approaches.



Methods

Supervised machine learning methods achieve good performance, but **unsupervised methods** are being introduced because many software defect datasets often:

- lack historical data
- have unlabeled or incomplete data

Hence, two clustering techniques are studied in this project and compared to common supervised classifiers:

- **Naive Bayes**
- **Random Forest**
- **Neural Gas Clustering**
- **Spectral Clustering**

Performance of methods

Each supervised and unsupervised method, along with the four different feature selection approaches, is compared to the others by the following evaluation measures:

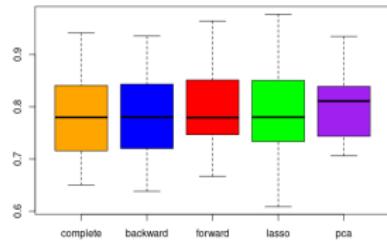
- **Accuracy** = $\frac{TP+TN}{TP+TN+FP+FN}$.
- Precision = $\frac{TP}{TP+FP}$.
- Recall = $\frac{TP}{TP+FN}$.
- Area Under the Curve (**AUC**): this curve is made by plotting the recall (or TPR, true positive rate) against the FPR (false positive rate), which can be computed as $1 - specificity$.

The comparison is focused in particular on Accuracy and AUC measures, the most common in software defect prediction.

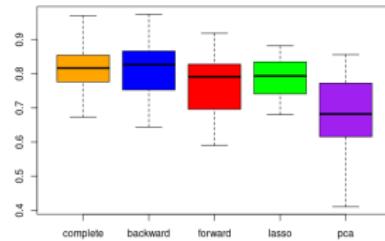
Performance of methods: accuracy

- According to accuracy, **Random Forest** achieves the **best performance** for all feature selection approaches.
- For **unsupervised** methods, **PCA** gives results **significantly worse** than the other feature selection approaches.

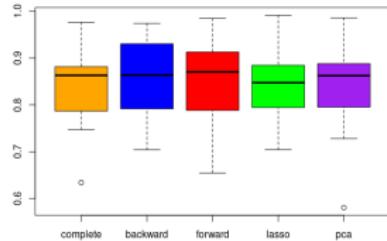
Accuracy - Naive Bayes



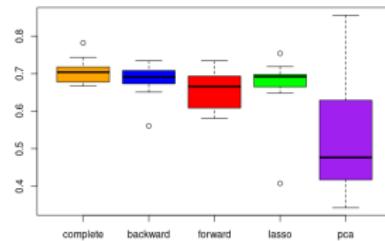
Accuracy - Neural Gas Clustering



Accuracy - Random Forest

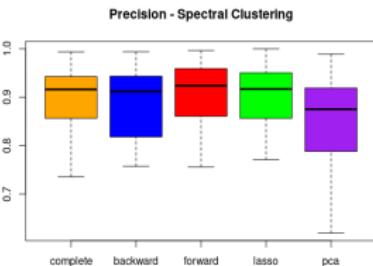
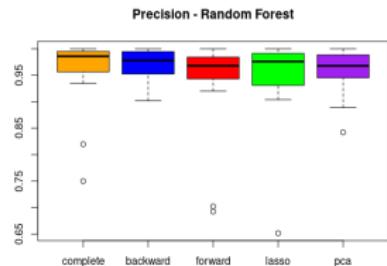
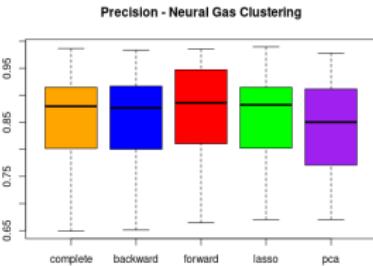
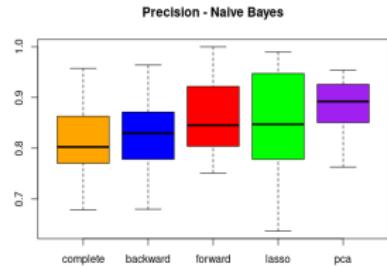


Accuracy - Spectral Clustering



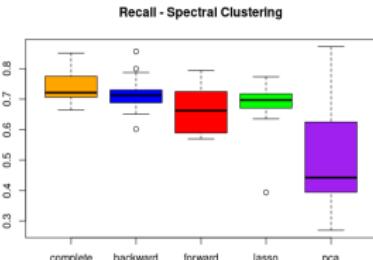
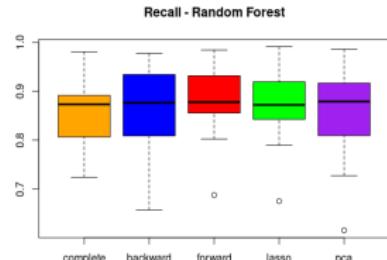
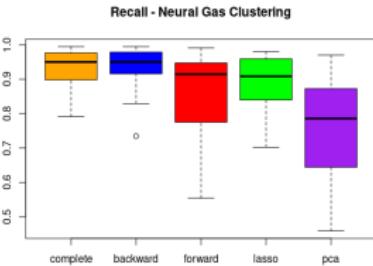
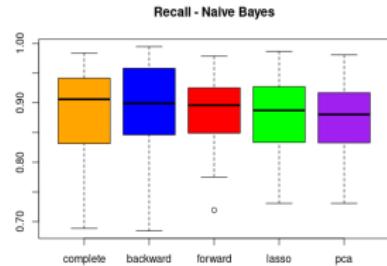
Performance of methods: precision

- **Precision** values are high (on average **above 0.8**) for all methods, and **Random Forest** achieves again the best performance.
- There is not a significant difference between different feature selection approaches. **PCA** performs as well as the others.



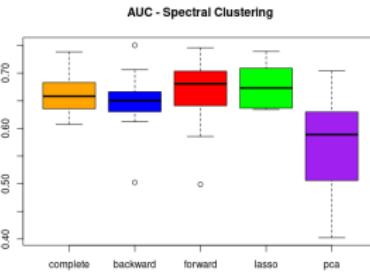
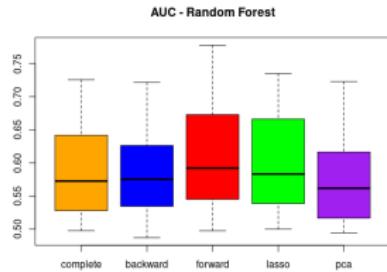
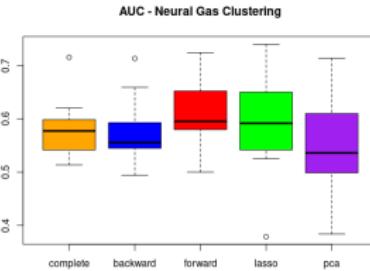
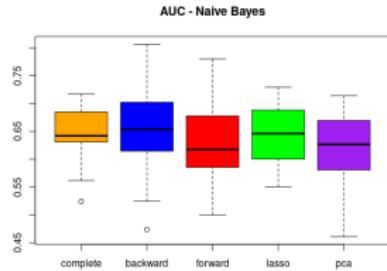
Performance of methods: recall

- The **recall** values are very high (**around 0.9**) for all the methods except spectral clustering. This means most non-defective classes are correctly identified.
- As for accuracy, **PCA negatively influences** the performance of **unsupervised** methods.



Performance of methods: AUC

- AUC values are lower than other evaluation measures, with **Spectral Clustering** having on average the best performance (**around 0.65**).
- For unsupervised methods, **PCA** performs worse than other feature selection approaches.



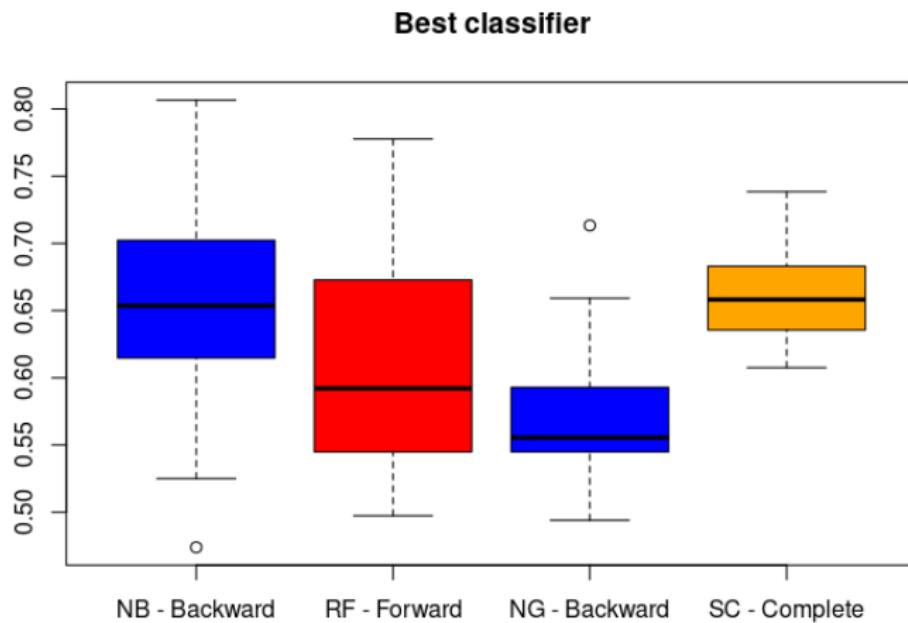
Comparison of classifiers

Every method of classification and clustering is compared with the different feature selections approaches, using:

- **Friedman test:** nonparametric test to compare if different methods have a significantly different performance. Tests the null hypothesis that different methods have the same performance.
- **Nemenyi test for multiple comparisons:** this nonparametric test is applied if the Friedman test rejects the null hypothesis.

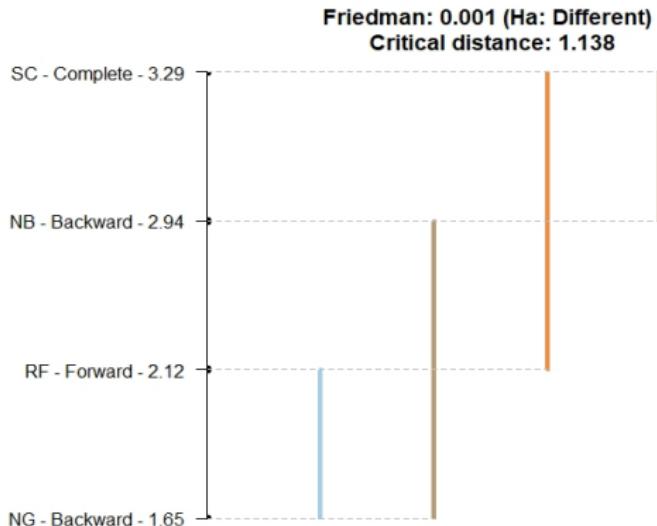
Comparison of classifiers: AUC

Naive Bayes and **Spectral Clustering** have the best AUC performance.
Random Forest achieves slightly lower results, but has higher variability.
Neural Gas Clustering has the lowest AUC values.



Comparison of classifiers: AUC

Spectral Clustering and **Naive Bayes** rankings do not differ by the Nemenyi critical distance. **Random Forest** ranking barely exceeds the critical distance and so performs as well as **Naive Bayes**, but not as well as **Spectral Clustering**. **Neural Gas Clustering** exceeds by far the critical distance: it achieves the worst performance.



Conclusion

RQ1. Do current unsupervised methods perform better or worse than the supervised ones?

- **Spectral Clustering** achieves average **0.66 AUC** on the complete datasets.
- **Naive Bayes** achieves average 0.65 AUC with backward selection.
- **Random Forest** achieves average 0.61 AUC with forward selection.
- However, their performances are **not significantly different**. The only method performing **worse** than the others is **Neural Gas Clustering**.

Conclusion

RQ2. Which feature selection method is the best for supervised methods?

- **Backward** selection is the best for **Naive Bayes**.
- **Forward** selection is the best for **Random Forest**.
- However, different **feature selection approaches do not influence significantly the performance of supervised methods**.

RQ3. Which feature selection method is the best for unsupervised methods?

- Backward selection is the best for **Neural Gas Clustering**.
- Using the **complete datasets** is the best approach for **Spectral Clustering**.
- All the feature selection approaches perform similarly, except for **PCA** which performs the worst.

Thanks for your attention

For more information, see the publications:

- E. Ronchieri, M. Canaparo, G. Bertaccini - A framework to conduct feature selection on software metric datasets, December 2021, in Proceedings of SDPS 2021, Workshop on Smart Pervasive Computing, <https://sdpsnet.org/sdps/documents/sdps-2021/SDPS%202021%20Proceedings.pdf#nameddest=15>.
- E. Ronchieri, M. Canaparo, G. Bertaccini - Software defect prediction: A study on software metrics using statistical and machine learning methods, March 2022, presented at ISGC 2022, paper under submission.

R framework

R framework at <https://github.com/rsma-defect-prediction>.

The screenshot shows a list of GitHub repositories belonging to the user 'rsma'. The repositories are:

- rsma.documentation** (Public) - Updated 3 minutes ago
- rsma.eclipse** (Private) - Updated 2 days ago
- rsma.nasa** (Private) - Updated 2 days ago
- rsma** (Public) - R software metrics analysis. Updated 2 days ago.
- rsma.demo** (Private) - Updated 2 days ago. Contains an R icon.
- rsma.quality** (Private) - Updated 2 days ago.
- rsma.github** (Private) - Updated 2 days ago.

At the top right, there is a 'Top languages' section with 'R' selected. The bottom of the screen features standard browser navigation icons.

Appendix

Number of features selected by different feature selection approaches.

Dataset	Complete	Backward	Forward	LASSO	PCA
Eclipse	18	10	7	5	11
Equinox	18	7	6	9	10
Lucene	18	9	5	15	11
Mylyn	18	9	7	11	11
PDE	18	13	7	14	10
CM1	38	13	10	9	10
MW2	38	13	7	8	11
PC1	38	19	14	18	11
PC3	38	23	10	26	11
PC4	38	21	19	29	14
JM1	22	12	14	21	9
KC1	22	13	12	16	8
KC3	40	14	7	12	9
MC2	40	20	5	8	10
MC1	39	16	9	15	15
PC5	39	22	15	27	14
PC2	37	11	6	8	8